

# Postgres Runs:

## Thumb Rule-1Runs:

Creating table, adding values from CSVs, Creating Primary keys with cluster indexing:

```
advdb=# CREATE TABLE trades_uniform (
    stock_symbol VARCHAR(10),
    time BIGINT,
    quantity INT,
    price NUMERIC(10, 2)
);
CREATE TABLE
advdb=# CREATE TABLE trades_fractal (
    stock_symbol VARCHAR(10),
    time BIGINT,
    quantity INT,
    price NUMERIC(10, 2)
);
CREATE TABLE
advdb=# COPY trades_uniform (stock_symbol, time, quantity, price)
FROM '/home/sc10670/ADB/advDB-HW2/Q2/uniform_trades.csv'
DELIMITER ','
CSV HEADER;
COPY 10000000
advdb=# COPY trades_fractal (stock_symbol, time, quantity, price)
FROM '/home/sc10670/ADB/advDB-HW2/Q2/fractal_trades.csv'
DELIMITER ','
CSV HEADER;
COPY 10000000
advdb=# ALTER TABLE trades_uniform ADD PRIMARY KEY (stock_symbol, time);
ALTER TABLE
advdb=# ALTER TABLE trades_fractal ADD PRIMARY KEY (stock_symbol, time);
ALTER TABLE
advdb=# CLUSTER trades_uniform USING trades_uniform_pkey;
CLUSTER
advdb=# CLUSTER trades_fractal USING trades_fractal_pkey;
CLUSTER
```

Running Query on Uniform Dataset with Clustered Indexing:

```
advdb=# EXPLAIN ANALYZE
SELECT stock_symbol, time, quantity, price,
       SUM(quantity) AS total_quantity,
       AVG(price) AS avg_price
  FROM trades_uniform
 WHERE time BETWEEN 100000 AND 800000
   AND price > 100
   AND quantity BETWEEN 500 AND 8000
 GROUP BY stock_symbol, time, quantity, price
 ORDER BY total_quantity DESC, avg_price ASC;
                                         QUERY PLAN
-----
Sort  (cost=316714.13..317923.64 rows=483803 width=63) (actual time=10732.557..12411.067 rows=470037 loops=1)
  Sort Key: (sum(quantity)) DESC, (avg(price))
  Sort Method: external merge Disk: 25368kB
    -> Finalize GroupAggregate  (cost=190177.70..252840.31 rows=483803 width=63) (actual time=1702.504..8926.413 rows=470037 loops=1)
        Group Key: stock_symbol, "time"
        -> Gather Merge  (cost=190177.70..241753.15 rows=403170 width=63) (actual time=1702.430..5439.377 rows=470037 loops=1)
            Workers Planned: 2
            Workers Launched: 2
            -> Partial GroupAggregate  (cost=189177.68..194217.30 rows=201585 width=63) (actual time=1665.395..3297.630 rows=156679 loops=3)
                Group Key: stock_symbol, "time"
                -> Sort  (cost=189177.68..189681.64 rows=201585 width=23) (actual time=1665.329..2183.321 rows=156679 loops=3)
                    Sort Key: stock_symbol, "time"
                    Sort Method: external merge Disk: 5536kB
                    Worker 0: Sort Method: external merge Disk: 5288kB
                    Worker 1: Sort Method: external merge Disk: 5312kB
                    -> Parallel Seq Scan on trades_uniform  (cost=0.00..167280.00 rows=201585 width=23) (actual time=0.435..1191.903 rows=156679 loops=3)
                        Filter: ((("time" >= 100000) AND ("time" <= 800000) AND (price > '100'::numeric) AND (quantity >= 500) AND (quantity <= 8000))
                        Rows Removed by Filter: 3176654
[ Planning Time: 10.149 ms
  Execution Time: 13547.375 ms
(20 rows)
```

Running Query on Fractal Dataset with Clustered Indexing:

```

-----  

QUERY PLAN  

-----  

Sort (cost=309575.74..310728.91 rows=461267 width=63) (actual time=10715.727..12529.905 rows=468152 loops=1)
  Sort Key: (sum(quantity) DESC, (avg(price)))
  Sort Method: external merge Disk: 25264kB
    -> Finalize GroupAggregate (cost=189091.76..248835.49 rows=461267 width=63) (actual time=1703.825..8901.576 rows=468152 loops=1)
      Group Key: stock_symbol, "time"
      -> Gather Merge (cost=189091.76..238264.78 rows=384390 width=63) (actual time=1703.760..5446.401 rows=468152 loops=1)
        Workers Planned: 2
        Workers Launched: 2
          -> Partial GroupAggregate (cost=188091.74..192896.61 rows=192195 width=63) (actual time=1689.259..3321.099 rows=156051 loops=3)
            Group Key: stock_symbol, "time"
            -> Sort (cost=188091.74..188572.22 rows=192195 width=23) (actual time=1689.166..2210.687 rows=156051 loops=3)
              Sort Key: stock_symbol, "time"
              Sort Method: external merge Disk: 5600kB
              Worker 0: Sort Method: external merge Disk: 5376kB
              Worker 1: Sort Method: external merge Disk: 5104kB
              -> Parallel Seq Scan on trades_fractal (cost=0.00..167280.00 rows=192195 width=23) (actual time=0.343..1221.364 rows=156051 loops=3)
                Filter: (("time" >= 100000) AND ("time" <= 800000) AND (price > '100'::numeric) AND (quantity >= 500) AND (quantity <= 8000))
                Rows Removed by Filter: 3177283
[ Planning Time: 3.432 ms
  Execution Time: 13503.324 ms
(20 rows)

```

## Running the Query to get output:

```

advdb=# SELECT stock_symbol, time, quantity, price,
  SUM(quantity) AS total_quantity,
  AVG(price) AS avg_price
FROM trades_uniform
WHERE time BETWEEN 100000 AND 800000
  AND price > 100
  AND quantity BETWEEN 500 AND 8000
GROUP BY stock_symbol, time, quantity, price
ORDER BY total_quantity DESC, avg_price ASC;
stock_symbol | time | quantity | price | total_quantity | avg_price
-----+-----+-----+-----+-----+-----+
s44195 | 796008 | 8000 | 111.00 | 8000 | 111.000000000000000000
s68886 | 579948 | 8000 | 119.00 | 8000 | 119.000000000000000000
s48222 | 237846 | 8000 | 124.00 | 8000 | 124.000000000000000000
s39386 | 396811 | 8000 | 129.00 | 8000 | 129.000000000000000000
s26148 | 112851 | 8000 | 133.00 | 8000 | 133.000000000000000000
s66240 | 313299 | 8000 | 154.00 | 8000 | 154.000000000000000000
s58862 | 474786 | 8000 | 172.00 | 8000 | 172.000000000000000000
s6459 | 464132 | 8000 | 173.00 | 8000 | 173.000000000000000000
s55302 | 458893 | 8000 | 173.00 | 8000 | 173.000000000000000000
s46889 | 656298 | 8000 | 177.00 | 8000 | 177.000000000000000000
s12420 | 643729 | 8000 | 180.00 | 8000 | 180.000000000000000000
s35361 | 220169 | 8000 | 185.00 | 8000 | 185.000000000000000000
s272213 | 437222 | 8000 | 189.00 | 8000 | 189.000000000000000000
s29541 | 481724 | 8000 | 209.00 | 8000 | 209.000000000000000000
s4534 | 756678 | 8000 | 209.00 | 8000 | 209.000000000000000000
s20057 | 731377 | 8000 | 212.00 | 8000 | 212.000000000000000000
s30353 | 145290 | 8000 | 219.00 | 8000 | 219.000000000000000000
s239 | 418984 | 8000 | 225.00 | 8000 | 225.000000000000000000
s13282 | 491977 | 8000 | 227.00 | 8000 | 227.000000000000000000
s42102 | 794912 | 8000 | 230.00 | 8000 | 230.000000000000000000
s52457 | 272736 | 8000 | 232.00 | 8000 | 232.000000000000000000
s2567 | 679497 | 8000 | 239.00 | 8000 | 239.000000000000000000
s463 | 484394 | 8000 | 251.00 | 8000 | 251.000000000000000000
s64978 | 212335 | 8000 | 256.00 | 8000 | 256.000000000000000000
s26876 | 442859 | 8000 | 257.00 | 8000 | 257.000000000000000000
s50418 | 319715 | 8000 | 263.00 | 8000 | 263.000000000000000000
s41638 | 662383 | 8000 | 265.00 | 8000 | 265.000000000000000000
s36124 | 584909 | 8000 | 273.00 | 8000 | 273.000000000000000000
s3316 | 214929 | 8000 | 286.00 | 8000 | 286.000000000000000000
s40940 | 251135 | 8000 | 287.00 | 8000 | 287.000000000000000000
s65400 | 282232 | 8000 | 288.00 | 8000 | 288.000000000000000000
s63333 | 617459 | 8000 | 290.00 | 8000 | 290.000000000000000000
s69873 | 545721 | 8000 | 300.00 | 8000 | 300.000000000000000000
s47609 | 742565 | 8000 | 304.00 | 8000 | 304.000000000000000000
s16326 | 214891 | 8000 | 315.00 | 8000 | 315.000000000000000000
s21264 | 122999 | 8000 | 318.00 | 8000 | 318.000000000000000000
s1890 | 144939 | 8000 | 320.00 | 8000 | 320.000000000000000000
s114 | 375481 | 8000 | 321.00 | 8000 | 321.000000000000000000
s47332 | 770992 | 8000 | 328.00 | 8000 | 328.000000000000000000
s20305 | 349253 | 8000 | 329.00 | 8000 | 329.000000000000000000
s2366 | 478673 | 8000 | 332.00 | 8000 | 332.000000000000000000
s5540 | 740409 | 8000 | 341.00 | 8000 | 341.000000000000000000
s39698 | 776697 | 8000 | 347.00 | 8000 | 347.000000000000000000
s61161 | 583331 | 8000 | 354.00 | 8000 | 354.000000000000000000
s62226 | 425751 | 8000 | 356.00 | 8000 | 356.000000000000000000
s57635 | 132206 | 8000 | 364.00 | 8000 | 364.000000000000000000
s35364 | 175141 | 8000 | 366.00 | 8000 | 366.000000000000000000
s39848 | 136702 | 8000 | 368.00 | 8000 | 368.000000000000000000
s51210 | 518891 | 8000 | 386.00 | 8000 | 386.000000000000000000
s60139 | 324381 | 8000 | 389.00 | 8000 | 389.000000000000000000
s4759 | 525587 | 8000 | 391.00 | 8000 | 391.000000000000000000
s32393 | 730299 | 8000 | 394.00 | 8000 | 394.000000000000000000
s7037 | 190898 | 8000 | 410.00 | 8000 | 410.000000000000000000
s68492 | 742004 | 8000 | 412.00 | 8000 | 412.000000000000000000
s53666 | 249227 | 8000 | 412.00 | 8000 | 412.000000000000000000
s56293 | 190225 | 8000 | 418.00 | 8000 | 418.000000000000000000
s57431 | 216729 | 8000 | 428.00 | 8000 | 428.000000000000000000
s30814 | 729853 | 8000 | 438.00 | 8000 | 438.000000000000000000
s61722 | 195124 | 8000 | 444.00 | 8000 | 444.000000000000000000
s2476 | 493622 | 8000 | 454.00 | 8000 | 454.000000000000000000

```

## Run 2:

## Creating Non Clustered Indexing for the second run::

```
[advdb=# ALTER TABLE trades_uniform DROP CONSTRAINT trades_uniform_pkey;
ALTER TABLE
[advdb=# ALTER TABLE trades_fractal DROP CONSTRAINT trades_fractal_pkey;
ALTER TABLE
[advdb=# CREATE INDEX idx_uniform_nonclustered ON trades_uniform (stock_symbol, time);
CREATE INDEX
[advdb=# CREATE INDEX idx_fractal_nonclustered ON trades_fractal (stock_symbol, time);
CREATE INDEX
advdb=# ]]
```

## Run for fractal Dataset:

```
advdb# EXPLAIN ANALYZE
SELECT stock_symbol, time, quantity, price,
       SUM(quantity) AS total_quantity,
       AVG(price) AS avg_price
FROM trades_fractal
WHERE time BETWEEN 100000 AND 800000
  AND price > 100
  AND quantity BETWEEN 500 AND 8000
GROUP BY stock_symbol, time, quantity, price
ORDER BY total_quantity DESC, avg_price ASC;
                                         QUERY PLAN
-----
Sort  (cost=316478.58..317662.91 rows=473730 width=63) (actual time=11366.936..13050.872 rows=468152 loops=1)
  Sort Key: (sum(quantity)) DESC, (avg(price))
  Sort Method: external merge Disk: 25264kB
    -> Finalize GroupAggregate (cost=189698.46..254009.26 rows=473730 width=63) (actual time=2355.243..9558.719 rows=468152 loops=1)
      Group Key: stock_symbol, "time", quantity, price
      -> Gather Merge (cost=189690.46..241179.05 rows=394776 width=63) (actual time=2355.186..5963.823 rows=468152 loops=1)
        Workers Planned: 2
        Workers Launched: 2
          -> Partial GroupAggregate (cost=188698.44..194612.08 rows=197388 width=63) (actual time=2240.829..3951.350 rows=156051 loops=3)
            Group Key: stock_symbol, "time", quantity, price
            -> Sort (cost=188698.44..189183.91 rows=197388 width=23) (actual time=2240.769..2829.924 rows=156051 loops=3)
              Sort Key: stock_symbol, "time", quantity, price
              Sort Method: external merge Disk: 5376kB
              Worker 0: Sort Method: external merge Disk: 5376kB
              Worker 1: Sort Method: external merge Disk: 5328kB
              -> Parallel Seq Scan on trades_fractal (cost=0.00..167280.00 rows=197388 width=23) (actual time=7.677..1144.793 rows=156051 loops=3)
                Filter: (("time" >= 100000) AND ("time" <= 800000) AND (price > '100'::numeric) AND (quantity >= 500) AND (quantity <= 8000))
                Rows Removed by Filter: 3177283
Planning Time: 57.844 ms
Execution Time: 14237.571 ms
(20 rows)
```

## Run for Uniform Dataset:

```
advdb# EXPLAIN ANALYZE
SELECT stock_symbol, time, quantity, price,
       SUM(quantity) AS total_quantity,
       AVG(price) AS avg_price
FROM trades_uniform
WHERE time BETWEEN 100000 AND 800000
  AND price > 100
  AND quantity BETWEEN 500 AND 8000
GROUP BY stock_symbol, time, quantity, price
ORDER BY total_quantity DESC, avg_price ASC;
                                         QUERY PLAN
-----
Sort  (cost=317110.75..318299.96 rows=475685 width=63) (actual time=11644.086..13282.170 rows=470037 loops=1)
  Sort Key: (sum(quantity)) DESC, (avg(price))
  Sort Method: external merge Disk: 25368kB
    -> Finalize GroupAggregate (cost=189785.44..254369.49 rows=475685 width=63) (actual time=240.2413..9827.377 rows=470037 loops=1)
      Group Key: stock_symbol, "time", quantity, price
      -> Gather Merge (cost=189785.44..241486.36 rows=396404 width=63) (actual time=2402.370..6207.556 rows=470037 loops=1)
        Workers Planned: 2
        Workers Launched: 2
          -> Partial GroupAggregate (cost=188785.42..194731.48 rows=198202 width=63) (actual time=2248.428..3961.938 rows=156679 loops=3)
            Group Key: stock_symbol, "time", quantity, price
            -> Sort (cost=188785.42..189280.92 rows=198202 width=23) (actual time=2248.363..2835.504 rows=156679 loops=3)
              Sort Key: stock_symbol, "time", quantity, price
              Sort Method: external merge Disk: 5872kB
              Worker 0: Sort Method: external merge Disk: 5256kB
              Worker 1: Sort Method: external merge Disk: 5016kB
              -> Parallel Seq Scan on trades_uniform (cost=0.00..167280.00 rows=198202 width=23) (actual time=12.770..1136.336 rows=156679 loops=3)
                Filter: ((("time" >= 100000) AND ("time" <= 800000) AND (price > '100'::numeric) AND (quantity >= 500) AND (quantity <= 8000))
                Rows Removed by Filter: 3176654
Planning Time: 19.289 ms
Execution Time: 14270.868 ms
(20 rows)
```

# Thumb Rule #2 Runs:

Creating tables for Skewed and Uniform Data:

```
[sc10670@crunchy1 Q2]$ psql -p 34567 -h /home/sc10670/pgsql/data/run
psql (13.16)
Type "help" for help.

sc10670=# CREATE TABLE uniform_data (
    id BIGINT PRIMARY KEY,
    user_id BIGINT NOT NULL
);
CREATE TABLE
sc10670=# CREATE TABLE skewed_data (
    id BIGINT PRIMARY KEY,
    user_id BIGINT NOT NULL
);
CREATE TABLE
sc10670=# COPY uniform_data (id, user_id)
FROM '/home/sc10670/ADB/advDB-HW2/Q2/uniform_data_thumb2.csv'
DELIMITER ','
CSV HEADER;
COPY 10000000
[sc10670=# COPY skewed_data (id, user_id)
FROM '/home/sc10670/ADB/advDB-HW2/Q2/skewed_data_thumb2.csv'
DELIMITER ','
CSV HEADER;
COPY 10000000
[sc10670=# ]
```

Running query on Uniform Data without Indexing:

```
sc10670=# EXPLAIN ANALYZE
SELECT user_id, COUNT(*) AS frequency
FROM uniform_data
WHERE user_id BETWEEN 10000 AND 20000
GROUP BY user_id
ORDER BY frequency DESC;
                                         QUERY PLAN
-----
Sort  (cost=345080.72..346634.58 rows=621543 width=16) (actual time=8314.985..8535.101 rows=99997 loops=1)
  Sort Key: (count(*)) DESC
  Sort Method: external merge Disk: 2552kB
    -> HashAggregate  (cost=260582.01..274648.71 rows=621543 width=16) (actual time=7567.132..8051.232 rows=99997 loops=1)
        Group Key: user_id
        Planned Partitions: 16  Batches: 17  Memory Usage: 4241kB  Disk Usage: 30896kB
          -> Seq Scan on uniform_data  (cost=0.00..204052.90 rows=1004962 width=8) (actual time=91.263..4612.659 rows=999158 l
oops=1)
              Filter: ((user_id >= 10000) AND (user_id <= 20000))
              Rows Removed by Filter: 9000842
Planning Time: 19.286 ms
Execution Time: 9148.476 ms
(11 rows)

sc10670=# EXPLAIN ANALYZE
SELECT user_id, COUNT(*) AS frequency
FROM skewed_data
WHERE user_id BETWEEN 1 AND 1000
GROUP BY user_id
ORDER BY frequency DESC;
```

## Running query on Skewed Data without Indexing:

```
sc10670=# EXPLAIN ANALYZE
SELECT user_id, COUNT(*) AS frequency
FROM skewed_data
WHERE user_id BETWEEN 1 AND 1000
GROUP BY user_id
ORDER BY frequency DESC;
                                         QUERY PLAN
-----
Sort  (cost=138733.93..138736.31 rows=951 width=16) (actual time=19351.391..22198.332 rows=1000 loops=1)
  Sort Key: (count(*)) DESC
  Sort Method: quicksort Memory: 71kB
->  Finalize GroupAggregate (cost=138445.95..138686.89 rows=951 width=16) (actual time=19326.888..22193.614 rows=1000 loop
s=1)
    Group Key: user_id
    ->  Gather Merge (cost=138445.95..138667.87 rows=1902 width=16) (actual time=19326.848..22183.699 rows=3000 loops=1)
        Workers Planned: 2
        Workers Launched: 2
        ->  Sort (cost=137445.93..137448.30 rows=951 width=16) (actual time=19294.287..19296.723 rows=1000 loops=3)
            Sort Key: user_id
            Sort Method: quicksort Memory: 71kB
            Worker 0: Sort Method: quicksort Memory: 71kB
            Worker 1: Sort Method: quicksort Memory: 71kB
            ->  Partial HashAggregate (cost=137389.37..137398.88 rows=951 width=16) (actual time=19288.305..19291.09
6 rows=1000 loops=3)
                Group Key: user_id
                Batches: 1  Memory Usage: 193kB
                Worker 0: Batches: 1  Memory Usage: 193kB
                Worker 1: Batches: 1  Memory Usage: 193kB
                ->  Parallel Seq Scan on skewed_data (cost=0.00..116556.09 rows=4166656 width=8) (actual time=0.36
5..9621.322 rows=3333333 loops=3)
                    Filter: ((user_id >= 1) AND (user_id <= 1000))
Planning Time: 1.474 ms
Execution Time: 22200.418 ms
(22 rows)
```

## Running query on Uniform Data with Indexing:

```
sc10670=# CREATE INDEX idx_user_id ON uniform_data(user_id);
CREATE INDEX
sc10670=# CREATE INDEX idx_user_id2 ON skewed_data(user_id);
CREATE INDEX
sc10670=# EXPLAIN ANALYZE
SELECT user_id, COUNT(*) AS frequency
FROM uniform_data
WHERE user_id BETWEEN 100000 AND 200000
GROUP BY user_id
ORDER BY frequency DESC;
                                         QUERY PLAN
-----
Sort  (cost=186332.73..107886.60 rows=621547 width=16) (actual time=5229.736..5448.095 rows=99997 loops=1)
  Sort Key: (count(*)) DESC
  Sort Method: external merge Disk: 2552kB
->  GroupAggregate (cost=0.43..35900.31 rows=621547 width=16) (actual time=0.978..4964.843 rows=99997 loops=1)
    Group Key: user_id
    ->  Index Only Scan using idx_user_id on uniform_data (cost=0.43..24659.96 rows=1004976 width=8) (actual time=0.928..
.2431.084 rows=999158 loops=1)
        Index Cond: ((user_id >= 100000) AND (user_id <= 200000))
        Heap Fetches: 0
Planning Time: 2.438 ms
Execution Time: 5648.564 ms
(10 rows)

sc10670=# EXPLAIN ANALYZE
SELECT user_id, COUNT(*) AS frequency
FROM skewed_data
WHERE user_id BETWEEN 1 AND 1000
GROUP BY user_id
ORDER BY frequency DESC;
```

## Running query on Skewed Data with Indexing:

```
sc10670=# EXPLAIN ANALYZE
SELECT user_id, COUNT(*) AS frequency
FROM skewed_data
WHERE user_id BETWEEN 1 AND 1000
GROUP BY user_id
ORDER BY frequency DESC;
                                         QUERY PLAN
-----
Sort  (cost=138732.89..138735.27 rows=951 width=16) (actual time=19936.806..19939.001 rows=1000 loops=1)
  Sort Key: (count(*)) DESC
  Sort Method: quicksort Memory: 71kB
->  Finalize GroupAggregate (cost=138444.91..138685.85 rows=951 width=16) (actual time=19912.477..19934.356 rows=1000 loop
s=1)
    Group Key: user_id
    ->  Gather Merge (cost=138444.91..138666.83 rows=1902 width=16) (actual time=19912.433..19924.414 rows=3000 loops=1)
        Workers Planned: 2
        Workers Launched: 2
        ->  Sort (cost=137444.89..137447.27 rows=951 width=16) (actual time=19900.879..19903.455 rows=1000 loops=3)
            Sort Key: user_id
            Sort Method: quicksort Memory: 71kB
            Worker 0: Sort Method: quicksort Memory: 71kB
            Worker 1: Sort Method: quicksort Memory: 71kB
            ->  Partial HashAggregate (cost=137388.33..137397.85 rows=951 width=16) (actual time=19894.984..19897.77
2 rows=1000 loops=3)
                Group Key: user_id
                Batches: 1  Memory Usage: 193kB
                Worker 0: Batches: 1  Memory Usage: 193kB
                Worker 1: Batches: 1  Memory Usage: 193kB
                ->  Parallel Seq Scan on skewed_data (cost=0.00..116555.00 rows=4166667 width=8) (actual time=72.7
80..10166.368 rows=3333333 loops=3)
                    Filter: ((user_id >= 1) AND (user_id <= 1000))
Planning Time: 1.903 ms
Execution Time: 19941.123 ms
(22 rows)
```